

EE 583 Pattern Recognition HW5

Eda Özkaynar 2375582

QUESTION 1




Fig. 1: Dataset

As can be seen in 2, this dataset consists of two distinct classes. However, its structure resembles an XOR problem, making it impossible to classify the data accurately using linear classifiers due to the non-linear decision boundaries required to separate the classes.

The network structure is a fully connected neural network, with one input layer, one hidden layer, and one output layer. The learning rate determines the step size at which the optimizer updates the model's weights. A small learning rate ensures gradual convergence but can be slow, while a large learning rate may result in instability. The optimizer determines how the model's parameters (weights and biases) are updated based on the gradients computed during backpropagation. In this case, Stochastic Gradient Descent (SGD) is used as the optimizer. Instead of utilizing the entire dataset in each iteration, SGD selects a single random training example or a small batch to calculate the gradient and update the model parameters, making it computationally efficient for large datasets. The loss function, which in this case is Mean Squared Error (MSE), quantifies the error between the model's predictions and the ground truth. During training, the optimizer aims to minimize this loss function by iteratively updating the model's parameters.

Metric	Value
Validation Loss (After Training)	0.217
Validation Accuracy	61.00%
Epoch 100 Train Loss	0.208

TABLE I: Training and Validation Metrics




Fig. 2: Learning Curve for training data

QUESTION 2




Fig. 3: Learning curve for 500 epoch

Metric	Value
Validation Loss (After Training)	0.093
Validation Accuracy	96.00%
Epoch 500 Train Loss	0.092

TABLE II: Training and Validation Metrics (Epoch 500)

When we compare I and II, we can easily observe that increasing the number of epochs improves the accuracy and reduces the loss. This occurs because the model is allowed more iterations to learn the underlying patterns in the data, effectively minimizing the error between the predicted outputs and the ground truth.

QUESTION 3

A. Modified Fully Connected Network

- **Input Layer:**
 - Accepts input of size `input_size`.
- **First Hidden Layer:**
 - Fully connected layer with input size `input_size` and output size `hidden_size`.
 - Applies ReLU activation function.
- **Second Hidden Layer:**
 - Fully connected layer with input size `hidden_size` and output size $2 \times \text{hidden_size}$.
 - Applies ReLU activation function.
- **Third Hidden Layer:**
 - Fully connected layer with input size $2 \times \text{hidden_size}$ and output size $4 \times \text{hidden_size}$.
 - Applies ReLU activation function.
- **Output Layer:**
 - Fully connected layer with input size $4 \times \text{hidden_size}$ and output size `num_classes`.
 - No activation function is applied.




Fig. 4: Learning Curve of Modified FCN

Metric	Value
Validation Loss (After Training)	0.193
Validation Accuracy	85.00%
Epoch 500 Train Loss	0.182


TABLE III: Training and Validation Metrics for Modified Fully Connected Network

Comparing II and III, we see that increasing the complexity of the model results in higher training loss and lower accuracy. This happens for a few reasons.

First, more complex models, with additional layers or more parameters, need larger datasets to learn effectively. If the training data isn't enough, the model struggles to fit the data properly, which increases the training loss. More parameters also make optimization harder, as the model might get stuck in local minima or saddle points.

Second, the drop in accuracy is likely due to overfitting. Complex models can capture very detailed patterns, but this often includes noise or irrelevant details in the training data. As a result, the model doesn't generalize well to unseen data, leading to lower accuracy on validation or test sets.

QUESTION 4



(a) Learning curve for the First Fully Connected Network (Learning Rate = 0.01) (b) Learning curve for the Modified Fully Connected Network (Learning Rate = 0.01)

Fig. 5: Learning Curves (Learning Rate = 0.01)

The learning rate is a critical hyperparameter that determines the speed and stability of model convergence during training. A higher learning rate allows faster convergence but may cause instability or poor performance if the model diverges or overshoots the optimal point. Conversely, a lower learning rate results in more stable convergence but requires more time to reach an accurate solution and may get stuck in saddle points or local minima.

In our case, increasing the learning rate from 0.001 to 0.01 improved accuracy. This suggests that a learning rate of 0.001 was too small for the model to effectively converge within 500 epochs. The increase in learning rate also reduced the training loss for the modified FCN. However, the modified FCN exhibited a higher validation loss compared to the simpler FCN. This indicates that the modified FCN, being a more complex model, struggled to generalize and overfitted the training data slightly.

In summary, while increasing the learning rate improved both accuracy and training efficiency, it also highlighted the modified FCN's tendency to overfit due to its increased complexity.

Metric	Value
Validation Loss (After Training)	0.028
Validation Accuracy	100.00%
Epoch 500 Train Loss	0.040

TABLE IV: Training and Validation Metrics for First Fully Connected Model

Metric	Value
Validation Loss (After Training)	0.043
Validation Accuracy	97.00%
Epoch 500 Train Loss	0.030

TABLE V: Training and Validation Metrics for Modified Fully Connected Model

QUESTION 5




Fig. 6: Learning Curve (Dropout = 0.5)

Metric	Value
Epoch 500 Train Loss	0.173
Validation Loss (After Training)	0.101
Validation Accuracy	95.00%

TABLE VI: Training and Validation Metrics for FCN with dropout

Dropout is a regularization technique that improves the model's ability to generalize by preventing overfitting. It works by randomly "dropping out" (i.e., setting to zero) a fraction of the neurons during each forward and backward pass in training. This forces the network to rely on multiple independent representations of the data, rather than overfitting to specific patterns in the training set.

In our case, adding a dropout layer has effectively reduced overfitting, as evidenced by the validation error being lower than the training error. This indicates that the model is generalizing better to unseen data. However, dropout introduces some level of instability in the model's learning process, as the randomly dropped neurons can cause fluctuations in the loss during training. This is expected behavior and is a trade-off for improved generalization.

QUESTION 6

Momentum	Epoch 500 Train Loss	Validation Loss (After Training)	Validation Accuracy
0.50	0.075	0.074	98.00%
0.85	0.043	0.039	99.00%
0.90	0.041	0.038	99.00%
0.99	0.025	0.024	98.00%

TABLE VII: Training and Validation Metrics for Different Momentum Values

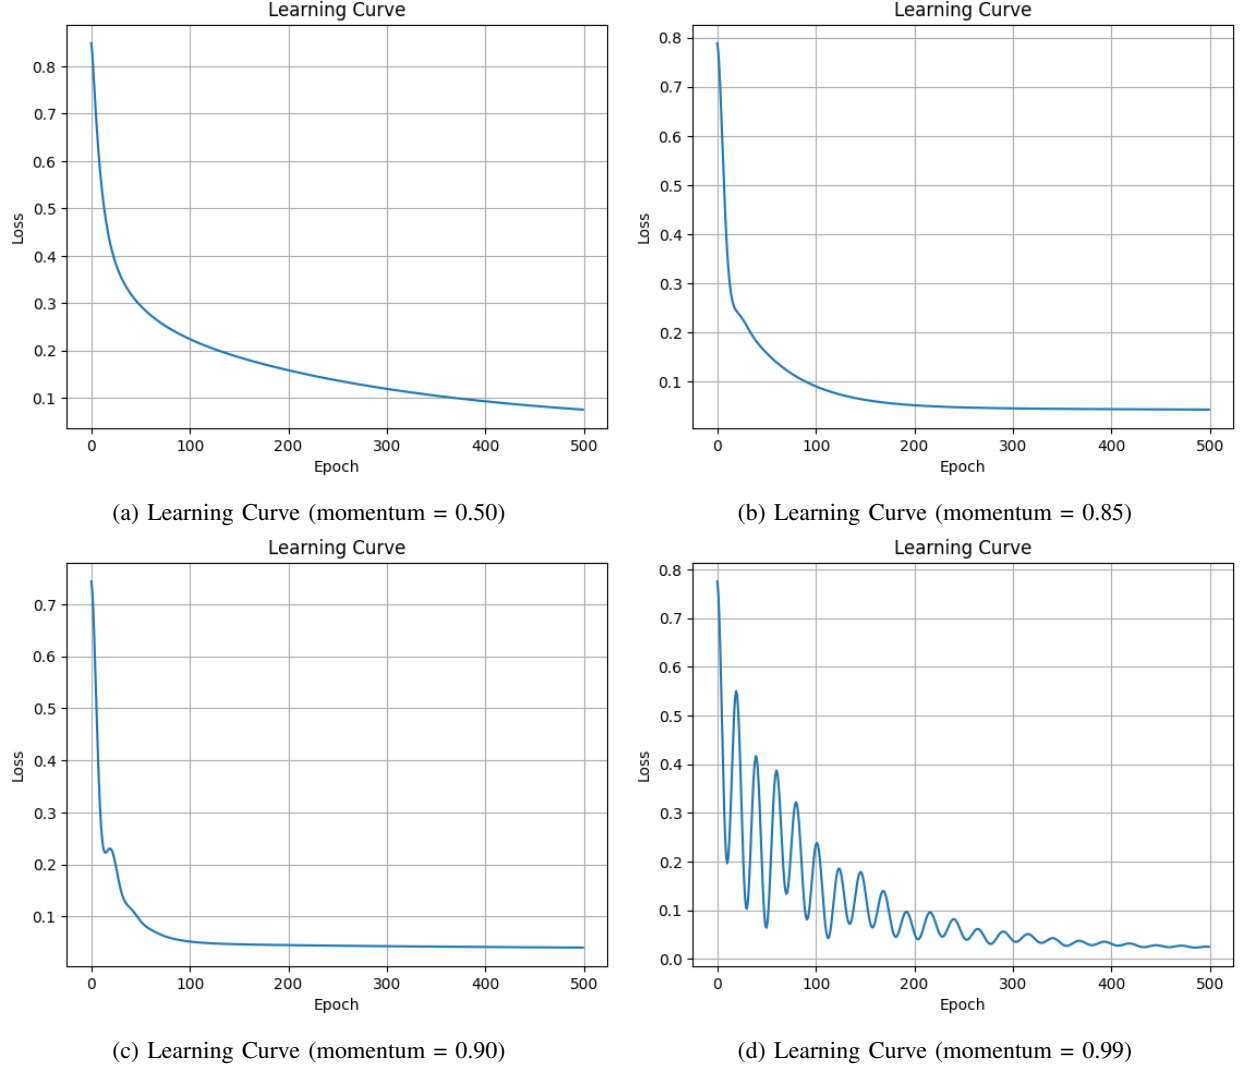


Fig. 7: Learning Curves for different momentum values

Momentum is a technique used in the Gradient Descent (GD) algorithm and is designed to speed up the optimization process and make it more stable. Momentum adds "acceleration" to parameter updates by taking into account past gradients. This allows optimization to be performed more effectively, especially in steep slope and oscillating regions. Momentum is typically used with $\gamma=0.9$ as the default and provides better results than standard Gradient Descent in most cases. In Figure 7, one can see that the model in 7c converges faster than models in 7a and 7b. However, in Figure 7d, the effect of high momentum causes the model to take larger steps with each iteration. This causes the loss function to initially fluctuate chaotically. Excessively high momentum value causes instability in the optimization process. Although the training loss reaches the lowest level, the verification accuracy is 98.00%, which is lower than other high momentum values.