

```
import torch
import numpy as np
import torchvision.datasets as datasets
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
from torchsummary import summary
```

```
#DEFINE YOUR DEVICE
```

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
```

```
print(device) #if cpu, go Runtime-> Change runtime type-> Hardware accelerator GPU -> Save -> Redo previous steps
```

```
→ cuda:0
```

```
#DOWNLOAD CIFAR-10 DATASET
```

```
train_data = datasets.CIFAR10('./data', train = True, download = True, transform = transforms.ToTensor())
```

```
test_data = datasets.CIFAR10('./data', train = False, transform = transforms.ToTensor())
```

```
→ Files already downloaded and verified
```

```
#DEFINE DATA GENERATOR
```

```
batch_size = 100
```

```
train_generator = torch.utils.data.DataLoader(train_data, batch_size = batch_size, shuffle = True)
```

```
test_generator = torch.utils.data.DataLoader(test_data, batch_size = batch_size, shuffle = False)
```

```
Üretilen kod lisansa tabi olabilir | Shefin-CSE16/Bangla-Handwritten-Character-Recognition
```

```
#DEFINE NEURAL NETWORK MODEL
```

```
class CNN(torch.nn.Module):
```

```
    def __init__(self):
```

```
        super(CNN, self).__init__()
```

```
        self.conv1 = torch.nn.Conv2d(3, 32, kernel_size = 3, stride = 1)
```

```
        self.conv2 = torch.nn.Conv2d(32, 64, kernel_size = 3, stride = 1)
```

```
        self.conv3 = torch.nn.Conv2d(64, 128, kernel_size = 3, stride = 1)
```

```
        self.conv4 = torch.nn.Conv2d(128, 256, kernel_size = 3, stride = 1)
```

```
        self.mpool = torch.nn.MaxPool2d(2)
```

```
        self.fc1 = torch.nn.Linear(1024, 2048) # Increased output size to 1024
```

```
        self.fc2 = torch.nn.Linear(2048, 256) # Increased output size to 256
```

```
        self.fc3 = torch.nn.Linear(256, 10)
```

```
        self.relu = torch.nn.ReLU()
```

```
        self.sigmoid = torch.nn.Sigmoid()
```

```
        self.drop = torch.nn.Dropout(0.4)
```

```
        self.bn1 = torch.nn.BatchNorm2d(32)
```

```
        self.bn2 = torch.nn.BatchNorm2d(64)
```

```
        self.bn3 = torch.nn.BatchNorm2d(128)
```

```
        self.bn4 = torch.nn.BatchNorm2d(256)
```

```
    def forward(self, x):
```

```
        hidden = self.bn1(self.relu(self.conv1(x)))
```

```
        hidden = self.mpool(self.bn2(self.relu(self.conv2(hidden))))
```

```
        hidden = self.mpool(self.bn3(self.relu(self.conv3(hidden))))
```

```
        hidden = self.mpool(self.bn4(self.relu(self.conv4(hidden))))
```

```
        hidden = hidden.view(-1,1024)
```

```
        hidden = self.relu(self.fc1(hidden))
```

```
        hidden = self.drop(hidden)
```

```
        hidden = self.relu(self.fc2(hidden))
```

```
        hidden = self.drop(hidden)
```

```
        output = self.fc3(hidden)
```

```
        return output
```

```
#CREATE MODEL
```

```
model = CNN()
```

```
model.to(device)
```

```
summary(model, (3,32,32))
```

```
→ -----
Layer (type)          Output Shape          Param #
=====
      Conv2d-1         [-1, 32, 30, 30]             896
      ReLU-2           [-1, 32, 30, 30]              0
  BatchNorm2d-3        [-1, 32, 30, 30]             64
      Conv2d-4         [-1, 64, 28, 28]          18,496
      ReLU-5           [-1, 64, 28, 28]              0
  BatchNorm2d-6        [-1, 64, 28, 28]          128
    MaxPool2d-7        [-1, 64, 14, 14]              0
      Conv2d-8         [-1, 128, 12, 12]         73,856
```

ReLU-9	[-1, 128, 12, 12]	0
BatchNorm2d-10	[-1, 128, 12, 12]	256
MaxPool2d-11	[-1, 128, 6, 6]	0
Conv2d-12	[-1, 256, 4, 4]	295,168
ReLU-13	[-1, 256, 4, 4]	0
BatchNorm2d-14	[-1, 256, 4, 4]	512
MaxPool2d-15	[-1, 256, 2, 2]	0
Linear-16	[-1, 2048]	2,099,200
ReLU-17	[-1, 2048]	0
Dropout-18	[-1, 2048]	0
Linear-19	[-1, 256]	524,544
ReLU-20	[-1, 256]	0
Dropout-21	[-1, 256]	0
Linear-22	[-1, 10]	2,570

=====

Total params: 3,015,690

Trainable params: 3,015,690

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 2.51

Params size (MB): 11.50

Estimated Total Size (MB): 14.03

#DEFINE LOSS FUNCTION AND OPTIMIZER

learning_rate = 0.001

loss_fun = torch.nn.CrossEntropyLoss()

optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate)

#TRAIN THE MODEL

model.train()

epoch = 25

num_of_batch=np.int32(len(train_generator.dataset)/batch_size)

loss_values = np.zeros(epoch*num_of_batch)

for i in range(epoch):

 for batch_idx, (x_train, y_train) in enumerate(train_generator):

 x_train, y_train = x_train.to(device), y_train.to(device)

 optimizer.zero_grad()

 y_pred = model(x_train)

 loss = loss_fun(y_pred, y_train)

 loss_values[num_of_batch*i+batch_idx] = loss.item()

 loss.backward()

 optimizer.step()

 if (batch_idx+1) % batch_size == 0:

 print('Epoch: {}/{} [Batch: {}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
 i+1, epoch, (batch_idx+1) * len(x_train), len(train_generator.dataset),
 100. * (batch_idx+1) / len(train_generator), loss.item()))



```

Epoch: 21/25 [Batch: 20000/50000 (40%)] Loss: 0.055289
Epoch: 21/25 [Batch: 30000/50000 (60%)] Loss: 0.025615
Epoch: 21/25 [Batch: 40000/50000 (80%)] Loss: 0.048198
Epoch: 21/25 [Batch: 50000/50000 (100%)] Loss: 0.113723
Epoch: 22/25 [Batch: 10000/50000 (20%)] Loss: 0.002894
Epoch: 22/25 [Batch: 20000/50000 (40%)] Loss: 0.069682
Epoch: 22/25 [Batch: 30000/50000 (60%)] Loss: 0.020700
Epoch: 22/25 [Batch: 40000/50000 (80%)] Loss: 0.049101
Epoch: 22/25 [Batch: 50000/50000 (100%)] Loss: 0.059412
Epoch: 23/25 [Batch: 10000/50000 (20%)] Loss: 0.141085
Epoch: 23/25 [Batch: 20000/50000 (40%)] Loss: 0.027072
Epoch: 23/25 [Batch: 30000/50000 (60%)] Loss: 0.117825
Epoch: 23/25 [Batch: 40000/50000 (80%)] Loss: 0.061632
Epoch: 23/25 [Batch: 50000/50000 (100%)] Loss: 0.060520
Epoch: 24/25 [Batch: 10000/50000 (20%)] Loss: 0.072651
Epoch: 24/25 [Batch: 20000/50000 (40%)] Loss: 0.048862
Epoch: 24/25 [Batch: 30000/50000 (60%)] Loss: 0.098408
Epoch: 24/25 [Batch: 40000/50000 (80%)] Loss: 0.243279
Epoch: 24/25 [Batch: 50000/50000 (100%)] Loss: 0.083427
Epoch: 25/25 [Batch: 10000/50000 (20%)] Loss: 0.039074
Epoch: 25/25 [Batch: 20000/50000 (40%)] Loss: 0.005645
Epoch: 25/25 [Batch: 30000/50000 (60%)] Loss: 0.136557
Epoch: 25/25 [Batch: 40000/50000 (80%)] Loss: 0.101074
Epoch: 25/25 [Batch: 50000/50000 (100%)] Loss: 0.135330


```

```
#PLOT THE LEARNING CURVE
```

```

iterations = np.linspace(0,epoch,num_of_batch*epoch)
plt.plot(iterations, loss_values)
plt.title('Learning Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid('on')

```



```
#TEST THE MODEL
```

```

model.eval()
correct=0
total=0

```

```

for x_val, y_val in test_generator:
    x_val = x_val.to(device)
    y_val = y_val.to(device)

```

```

output = model(x_val)
y_pred = output.argmax(dim=1)

```

```

for i in range(y_pred.shape[0]):
    if y_val[i]==y_pred[i]:
        correct += 1
    total +=1

```

```
print('Validation accuracy: %.2f%%' %((100*correct)/(total)))
```



```
Validation accuracy: 80.00%
```

